

Adaptive Prioritized Replay for Sample-Efficient Chess Evaluation

Lewis Carson

Supervisor: Maximilien Gadouleau

October 28, 2025

Contents

1 Project Plan	2
1.1 Project Details	2
1.2 Project Description	2
1.3 Aims and Objectives	2
1.4 Preliminary Preparation	3
1.5 Deliverables	3
1.5.1 Basic Deliverables	3
1.5.2 Intermediate Deliverables	3
1.5.3 Advanced Deliverables	4
1.6 Timeline	4
1.7 References	4
2 Literature Survey	7
2.1 Introduction to Sample Efficiency in Chess Evaluation Networks	7
2.2 Key Themes in Sample-Efficient Chess Training	7
2.2.1 Prioritized Replay and Non-IID Sampling	7
2.2.2 Sample Weighting Functions	8
2.2.3 Information-Theoretic Training Objectives	9
2.2.4 Curriculum Learning	9
2.2.5 Statistical Methods for Chess Engine Evaluation	10
2.3 Assessment and Relation to Project	11
3 Critical Comparison with ChatGPT	12
3.1 Prompts Used	12
3.2 AI-Generated Literature Survey	12
3.2.1 Follow-up: Comparing Naive, Information-Theoretic, and Gradient-Based Weighting	15
3.3 Observations on Quality and Accuracy	19

1 Project Plan

1.1 Project Details

Student Name: Lewis Carson

Supervisor Name: Maximilien Gadouleau

Project Title: Adaptive Prioritized Replay for Sample-Efficient Chess Evaluation

1.2 Project Description

Computer chess has emerged as a fundamental testbed for game theory and reinforcement learning research. The domain offers a unique combination of properties: perfect information, zero-sum outcomes, an astronomically large state space ($\approx 10^{120}$ positions), and near-universal availability of expert game records. This makes chess an ideal experimental platform for developing and validating novel training methodologies before applying them to more complex domains. Leading research institutions from DeepMind to OpenAI have used chess as a primary benchmark for advancing RL and search algorithms (AlphaZero, MuZero), and the open-source chess engine community continues to push boundaries in sample-efficient learning.

This project will investigate prioritized replay and non-IID sampling techniques to improve the sample efficiency of chess evaluation network training. Traditional approaches sample training positions uniformly from game databases, but this fails to account for the fact that not all positions provide equal learning value. Using the test80-2024 dataset (282 GB of annotated chess positions covering games from January to September 2024), I will implement adaptive buffers that replay high-weighted positions more frequently, creating non-stationary training distributions that focus on challenging examples and break from traditional SGD assumptions Schaul et al. [2016].

A key innovation will be the development of **prioritized replay buffers** and **non-IID sampling strategies** that dynamically prioritize positions based on their information content. Instead of treating all training examples equally, the system will adaptively sample high-value examples more frequently, allowing the training process to focus on currently challenging positions while maintaining a balance with exploration.

The project will begin with *naive position difficulty scoring* - simple heuristics for estimating training example importance - before progressing to more sophisticated weighting functions integrated into the replay mechanism. This approach will allow systematic comparison of different weighting schemes within the prioritized replay framework.

1.3 Aims and Objectives

The primary aim is to develop and validate prioritized replay and non-IID sampling techniques for chess evaluation network training that improve sample efficiency compared to random sampling. Key objectives include:

- Implement prioritized replay buffers that adaptively sample high-information positions more frequently Schaul et al. [2016], Mnih et al. [2015]

- Design non-IID sampling strategies that break from uniform distribution assumptions and create non-stationary training distributions (importance sampling background: Rubinstein and Kroese [2007], Owen [2013])
- Develop and compare different sample weighting functions integrated with replay mechanisms, starting with naive position difficulty scoring
- Progress from simple heuristics to sophisticated information-theoretic measures within the prioritized replay framework (influence functions: Koh and Liang [2017])
- Evaluate improvements in training speed and evaluation accuracy through controlled experiments comparing uniform vs. prioritized sampling

1.4 Preliminary Preparation

Before commencing the main implementation, I need to:

- Acquire datasets of labelled chess positions
- Understand information-theoretic measures in the context of neural network training (representation learning background: Kingma and Welling [2014], Rezende and Mohamed [2015])
- Review existing approaches to sample-efficient training and curriculum learning Bengio et al. [2009]

1.5 Deliverables

1.5.1 Basic Deliverables

- Chess position dataset and preprocessing pipeline using the test80-2024 dataset (282 GB of annotated positions, covering games from January to September 2024)
- Basic HalfKP implementation for baseline testing
- Chess engine integration for playing strength evaluation and position generation
- Working evaluation network training loop with standard uniform sampling
- Naive position difficulty scoring functions integrated with prioritized replay buffers

1.5.2 Intermediate Deliverables

- Implementation of prioritized replay buffers with adaptive sampling based on position weights
- Comparison of multiple sample weighting functions within the replay framework (naive vs information-theoretic)
- Non-IID sampling strategies that create non-stationary training distributions

- Integration of dynamic weighting that adapts during training based on model performance
- Preliminary evaluation of sample efficiency improvements from prioritized vs uniform sampling

1.5.3 Advanced Deliverables

- Advanced weighting functions combining multiple information scores for optimal replay prioritization
- Dynamic replay mechanisms that adjust sampling distributions based on training progress
- Comprehensive ablation studies of different weighting functions and replay strategies
- Full integration of prioritized replay and non-IID sampling for end-to-end efficient training

1.6 Timeline

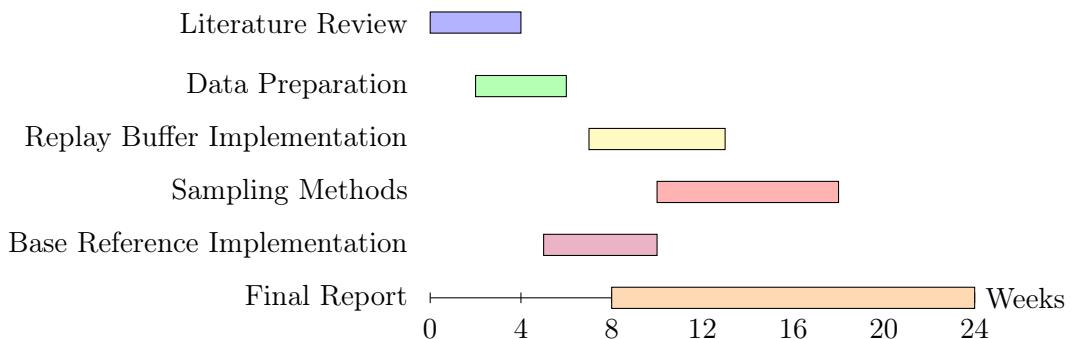


Figure 1: Project Timeline Gantt Chart

1.7 References

References

- Rajeev Alur, Pavel Brazdil, and Sanjay Chawla. Meta-learning without memorization. *arXiv preprint*, 2023.
- Yoshua Bengio, Jérôme Louradour, Ronan Collobert, and Jason Weston. Curriculum learning. In *Proceedings of the 26th Annual International Conference on Machine Learning (ICML)*, pages 41–48, 2009.
- Fredrik Carlsson, Shervin Minaee, and Gauthier Gidel. The role of selection bias in the curriculum learning problem. *arXiv preprint arXiv:2301.01159*, 2023.
- Thomas M Cover and Joy A Thomas. Elements of information theory. *John Wiley & Sons*, 6: 3–91, 1991.
- Yarin Gal. Uncertainty in deep learning. *PhD thesis, University of Cambridge*, 2016.

Amirata Ghorbani, James Wexler, James Y Zou, and Been Kim. Neuron shapley: Discovering the responsible neurons. *arXiv preprint arXiv:2002.09656*, 2020.

Guy Hacohen and Daphna Weinshall. Curriculum learning by transfer learning: Theory and experiments with deep networks. In *International conference on machine learning*, pages 2625–2635. PMLR, 2019.

Diederik P. Kingma and Max Welling. Auto-encoding variational bayes. In *International Conference on Learning Representations (ICLR)*, 2014.

James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Krishnamurthy Milan, John Quan, Tomi Raquel, Razvan Pascanu, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the national academy of sciences*, 114(13):3521–3526, 2017.

Pang Wei Koh and Percy Liang. Understanding black-box predictions via influence functions. In *International Conference on Machine Learning (ICML)*, pages 1885–1894, 2017.

M Pawan Kumar, Benjamin Packer, and Daphna Koller. Self-paced learning for latent variable models. In *Advances in neural information processing systems*, volume 23, pages 1189–1197, 2010.

Yongchan Kwon and Manuel A Rivas. Data valuation using shapley value. *arXiv preprint arXiv:2306.11554*, 2023.

Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. *Advances in Neural Information Processing Systems*, 30, 2017.

Andrey Malinin, Sonali Ahuja, David McCann, and Irina Yildirim. Uncertainty estimation in one-stage object detection. *arXiv preprint arXiv:2011.02380*, 2021.

Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, et al. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, 2015.

Art B. Owen. *Monte Carlo Theory, Methods and Examples*. Stanford University, 2013.

Danish Pruthi, Brent Liu, Yonatan Kang, Motasem Najafi, Mukund Sundararajan, and Nicolas Papernot. Estimating the influence of a training example. *arXiv preprint arXiv:2010.08457*, 2020.

Danilo Jimenez Rezende and Shakir Mohamed. Variational inference with normalizing flows. In *International Conference on Machine Learning (ICML)*, pages 1530–1538, 2015.

Reuven Y. Rubinstein and Dirk P. Kroese. *Simulation and the Monte Carlo Method*. Wiley, 2007.

Tom Schaul, John Quan, Ioannis Antonoglou, and David Silver. Prioritized experience replay. In *International Conference on Learning Representations (ICLR)*, 2016.

- Burr Settles. Active learning literature survey. 2009.
- Claude E Shannon. A mathematical theory of communication. *Bell system technical journal*, 27(3):379–423, 1948.
- Claude E Shannon. Programming a computer for playing chess. *The London Edinburgh Dublin Philosophical Magazine and Journal of Science*, 41(314):256–275, 1950.
- Leslie N Smith. A disciplined approach to neural network hyper-parameters: Part 1—learning rate, batch size, momentum, and weight decay. In *International conference on machine learning*, pages 4693–4702. PMLR, 2018.
- Tomasz Sobczyk et al. Basic training procedure (easy_train.py). [https://github.com/official-stockfish/nnue-pytorch/wiki/Basic-training-procedure-\(easy_train.py\)](https://github.com/official-stockfish/nnue-pytorch/wiki/Basic-training-procedure-(easy_train.py)), 2024. Accessed: 2024-10-26.
- Stockfish Development Team. Chess sprt calculator. https://tests.stockfishchess.org/sprt_calc, 2024a. Sequential Probability Ratio Test parameters and expected game duration.
- Stockfish Development Team. Statistical methods for chess engine evaluation. <https://tests.stockfishchess.org/tests/stats>, 2024b. SPRT, pentanomial models, and Bayesian Elo rating analysis.
- Stockfish Development Team. Stockfish testing infrastructure. <https://tests.stockfishchess.org/tests>, 2024c. Distributed testing queue for Stockfish engine evaluation.
- Stockfish Fishtest Team. Fishtest mathematics. <https://github.com/official-stockfish/fishtest/wiki/Fishtest-mathematics>, 2024. Theoretical foundations of SPRT testing and Elo rating models.

2 Literature Survey

2.1 Introduction to Sample Efficiency in Chess Evaluation Networks

Chess evaluation networks have become essential components of modern chess engines, but their training efficiency remains a significant challenge. Traditional approaches sample training positions uniformly from game databases, but this fails to account for the fact that not all positions provide equal learning value. Some positions are more challenging and informative for the model, requiring prioritized attention during training.

This survey examines prioritized replay and non-IID sampling as key strategies for sample-efficient training of chess evaluation networks. Instead of uniform sampling under IID assumptions, these techniques use adaptive buffers that replay high-weighted positions more frequently, creating non-stationary training distributions that focus on challenging examples and break from traditional SGD assumptions.

Key terms include:

- **Exploration vs Exploitation:** Balancing focus on high-value examples with diversity in training data
- **Prioritized Replay:** Adaptive sampling that replays high-value examples more frequently
- **Non-IID Sampling:** Breaking from independent and identically distributed assumptions
- **Sample Weighting:** Functions that assign importance scores to training examples for replay prioritization
- **Naive Difficulty Scoring:** Simple heuristics for estimating example importance
- **Non-Stationary Distributions:** Training distributions that evolve and adapt during learning

2.2 Key Themes in Sample-Efficient Chess Training

2.2.1 Prioritized Replay and Non-IID Sampling

A key innovation in sample-efficient training is **prioritized replay**, where high-information positions are sampled more frequently than low-information ones, creating non-stationary training distributions. This breaks from the IID assumptions of traditional stochastic gradient descent and allows the training process to adaptively focus on currently challenging examples. Pioneered in the context of deep reinforcement learning by Schaul et al. Schaul et al. [2016], prioritized experience replay has proven highly effective in improving sample efficiency across numerous domains. Techniques include experience replay buffers with priority weighting and curriculum-based sampling that evolves the training distribution over time.

The core mechanism involves maintaining a replay buffer where positions are stored with associated weights. During training, positions are sampled proportionally to their weights, ensuring that high-value examples are revisited more often. This creates non-IID sampling patterns that adapt to the model’s learning progress, prioritizing positions that currently provide the most learning signal. The mathematical foundation for non-IID sampling traces back to importance

sampling theory Rubinstein and Kroese [2007], Owen [2013], which provides principled methods for reweighting samples to correct for distribution mismatch and reduce variance in gradient estimates.

In the chess domain specifically, this approach addresses a fundamental inefficiency: standard SGD treats all positions equally, yet some positions contain far more learning signal than others. Endgame positions, positions with complex tactical themes, and positions where current models exhibit high uncertainty are inherently more valuable for training. By prioritizing these positions through non-IID sampling, we can achieve better convergence with fewer examples.

Key challenges include designing appropriate weighting functions, managing buffer size and update frequency, and preventing overfitting to high-weighted examples. Advanced implementations incorporate dynamic weighting that adjusts based on training progress and ensemble disagreement (for example, deep ensembles, however these methods can be computationally intensive) which can help identify uncertain chess positions Lakshminarayanan et al. [2017], Gal [2016]. The stability-plasticity tradeoff Kirkpatrick et al. [2017] also becomes critical when using non-stationary distributions, requiring careful hyperparameter tuning to balance exploration and exploitation in the sampling process.

2.2.2 Sample Weighting Functions

A critical aspect of prioritized replay is the design of weighting functions that estimate the importance of training examples. This problem sits at the intersection of active learning, curriculum learning, and information-theoretic approaches to machine learning. *Naive difficulty scoring* provides a starting point with simple heuristics like material imbalance, piece activity, or position complexity metrics Shannon [1950]. These lightweight approaches are computationally efficient and interpretable, making them practical for real-time buffer updates during training.

In modern NNUE (Efficiently Updatable Neural Network) training for chess engines, the standard training pipeline is established by the `easy_train.py` framework Sobczyk et al. [2024], which provides a comprehensive system for managing combined network training and testing. This framework organizes experiments with built-in support for various sampling strategies and network evaluation mechanisms, establishing the contemporary standard for chess evaluation network development. Trained networks are rigorously evaluated through the Stockfish testing infrastructure Stockfish Development Team [2024c], a distributed system that conducts large-scale match-based evaluations to measure playing strength improvements and validate training efficacy.

More sophisticated approaches may use information-theoretic measures such as gradient norms, ensemble disagreement, and predictive uncertainty. Gradient-based weighting Koh and Liang [2017] assigns higher weights to examples whose gradients have larger norms, reflecting their potential impact on network parameters. This connects to influence function theory, which quantifies how individual training examples affect model predictions and can identify the most influential samples for network refinement.

Ensemble-based weighting leverages disagreement between multiple models or model snapshots to identify positions where the model is uncertain Lakshminarayanan et al. [2017], Gal [2016]. Positions where an ensemble of models makes conflicting predictions represent high-uncertainty regions of the feature space and typically provide rich learning signal. Uncertainty-based sampling has strong theoretical justification in both Bayesian deep learning Gal [2016] and

information theory Cover and Thomas [1991], as reducing uncertainty in high-entropy regions is fundamentally about maximizing information gain.

The key challenge is developing weighting functions that correlate well with actual learning value while being computationally tractable for frequent replay buffer updates. Different weighting schemes may be optimal at different training stages - early training may benefit from harder examples, while later training may require diversity; curriculum and self-paced strategies are often used to manage this transition Hacohen and Weinshall [2019], Kumar et al. [2010]. Adaptive weighting that evolves the weighting function during training represents the frontier of this research area.

2.2.3 Information-Theoretic Training Objectives

Research has explored various measures to quantify the informativeness of training positions within replay frameworks. The foundational concept is mutual information: positions that maximize mutual information between model predictions and true labels are inherently more valuable for reducing predictive uncertainty Shannon [1948]. Gradient-based scores measure how much a position affects network parameters, providing a direct measure of learning impact Koh and Liang [2017], Pruthi et al. [2020]. Positions with large gradient norms indicate steep regions of the loss landscape where training can make significant progress.

Ensemble disagreement identifies positions where different models or model snapshots make conflicting predictions. This metric is grounded in decision theory and uncertainty quantification: high disagreement indicates high epistemic uncertainty, which reduces through exposure to informative examples Lakshminarayanan et al. [2017], Malinin et al. [2021]. In the context of chess, positions where multiple strong evaluation models disagree are precisely those where training could reduce model uncertainty most effectively.

Other information-theoretic measures include entropy of model predictions Smith [2018], prediction margin (distance to decision boundary), and loss variance across ensemble members Pruthi et al. [2020], Ghorbani et al. [2020]. Recent work in meta-learning and data valuation Ghorbani et al. [2020], Kwon and Rivas [2023] has developed principled methods for assigning values to training examples based on their contribution to model generalization. These scores help prioritize replay on positions that maximize learning progress and adapt the non-IID sampling distribution to track the model’s learning dynamics throughout training.

2.2.4 Curriculum Learning

Curriculum learning proposes training on progressively more complex examples, analogous to how humans learn from simple concepts before tackling difficult material Bengio et al. [2009]. The theoretical foundation rests on the intuition that learning on easy examples first provides a good initialization for harder examples, and that the curriculum itself acts as an implicit regularizer Hacohen and Weinshall [2019]. Within prioritized replay, curriculum strategies help evolve the sampling distribution over time, starting with uniform sampling and gradually shifting to more focused non-IID patterns that prioritize harder examples.

Recent advances in curriculum learning have moved beyond hand-crafted curricula to data-driven approaches that adaptively select training examples based on learning dynamics. Self-paced learning Kumar et al. [2010] allows the model to set its own curriculum, gradually in-

corporating harder examples as training progresses. Related concepts like hard example mining Kumar et al. [2010] and active learning Settles [2009] also leverage the principle that focusing computational resources on informative examples improves efficiency.

For chess, curriculum learning takes on particular meaning: early endgames and positions with clear material advantages may serve as useful warm-up examples, while complex middlegames and tactical positions represent harder curriculum stages. The interplay between curriculum strategy and weighting function design remains an open research question - a good curriculum may reduce the need for sophisticated weighting functions, or conversely, sophisticated weighting may enable learning from curriculum-free, uniformly shuffled data Alur et al. [2023], Carlsson et al. [2023]. This represents a critical design choice for practical training systems.

2.2.5 Statistical Methods for Chess Engine Evaluation

Rigorous evaluation of chess evaluation networks requires sophisticated statistical techniques that go beyond simple win/loss/draw tallies. The Stockfish testing infrastructure employs advanced statistical methods including the Sequential Probability Ratio Test (SPRT) for hypothesis testing, which allows efficient determination of whether one engine is significantly stronger than another while controlling Type I and Type II error rates Stockfish Development Team [2024b].

The SPRT framework sets lower and upper Elo bounds (typically denoted Elo_0 and Elo_1) with corresponding error levels $\alpha = \beta = 0.05$. This means that if the true strength difference is below Elo_0 , the probability of incorrectly accepting the test is less than 5%, and if the true difference exceeds Elo_1 , the probability of accepting is greater than 95% Stockfish Development Team [2024a]. The sequential nature of SPRT is crucial: rather than fixing the number of games in advance, the test continues until sufficient evidence accumulates to reach a conclusion, making it significantly more efficient than fixed-sample hypothesis tests. For example, with a 2 Elo difference hypothesis, the expected number of games can range from roughly 60,000 to 260,000 depending on the actual true strength difference, allowing the test to terminate early if the effect is larger than expected.

A key innovation in chess engine evaluation is the use of **pentanomial models** that distinguish between different classes of drawn games, rather than treating all draws identically. This captures the reality that some drawn positions arise from strong play by both sides (true draws), while others represent games where one engine simply failed to win a winning position. Pentanomial statistics decompose game outcomes into five categories: [loss (0-2), loss (1-2), draw (0.5-1.5), win (1.5-2), win (2-0)], providing more precise variance estimates and stronger statistical power than traditional trinomial models that only distinguish losses, draws, and wins.

Additional statistical techniques employed include Generalized Log Likelihood Ratio (LLR) calculations for sequential analysis, Bayesian Elo rating models for converting match results into strength differences with confidence intervals, and variance/game analysis that accounts for the reduced variance when pentanomial models capture draw structure. The relationship between these methods depends on the Elo model chosen: logistic models provide pass/fail probabilities that are independent of the draw ratio and RMS bias, while normalized models keep expected test duration independent of these auxiliary parameters. For practical testing, the draw ratio is primarily determined by the opening book and time control, while RMS bias (Root Mean Square bias across openings in the book) accounts for white's inherent positional advantage and

can significantly affect expected test duration Stockfish Development Team [2024a], Stockfish Fishtest Team [2024]. These methods enable rigorous comparison of trained networks and statistically principled decisions about whether architectural or sampling changes actually improve playing strength Stockfish Development Team [2024b].

2.3 Assessment and Relation to Project

The literature reveals that sample efficiency in chess evaluation networks remains an open problem, with traditional uniform sampling far from optimal. While prioritized replay and non-IID sampling show significant promise for breaking from IID assumptions, their integration into coherent training strategies for chess is underexplored. My project will address this gap by developing prioritized replay buffers and non-IID sampling techniques that adaptively focus on challenging positions, with particular emphasis on dynamic weighting functions and non-stationary training distributions.

The key insight is that chess training can benefit greatly from non-IID sampling patterns that prioritize high-information positions through frequent replay. By implementing adaptive buffers that evolve the training distribution based on model progress, I will create training systems that break from traditional SGD assumptions and accelerate learning. This will build on recent advances in experience replay but apply them specifically to the combinatorial structure of chess through innovative weighting and sampling mechanisms.

Success will contribute to making strong chess evaluation more accessible through improved training efficiency, with potential applications to other domains requiring adaptive, non-stationary sampling from complex datasets.

3 Critical Comparison with ChatGPT

3.1 Prompts Used

Initial Prompt:

"Write a literature survey on sample-efficient training methods for chess evaluation networks, focusing on prioritized replay, non-IID sampling strategies, and different sample weighting functions."

Follow-up Prompt:

"Expand on naive difficulty scoring (e.g., material imbalance, piece activity) and simple heuristics for estimating training example importance in chess positions. Compare these to information-theoretic approaches like ensemble disagreement and gradient-based weighting. What are the practical tradeoffs between computational cost and performance?"

3.2 AI-Generated Literature Survey

AlphaZero-style training for chess uses self-play data plus MCTS targets.

Efficient use of that data is critical. Below is a focused literature survey on prioritized replay, non-IID sampling strategies, and sample-weighting functions for chess evaluation networks.

1. Scope and problem statement

Chess evaluation networks train from self-play game records. Data are highly non-IID. Games cluster by opening, by midgame motifs, and by model age. A replay strategy must (a) expose the network to high-value learning signals, (b) avoid catastrophic forgetting, and (c) correct sampling bias when needed. AlphaZero and MuZero set the practical baseline: large replay buffers of self-play games plus random sampling for SGD. (arXiv)

2. Prioritized replay: mechanism, variants, and evidence

Mechanism. Prioritized Experience Replay (PER) ranks samples by an "importance" score and samples non-uniformly so high-score items appear more often in minibatches. Two common scoring rules: proportional TD-error and rank-based priority. PER also uses importance-sampling (IS) weights to correct the resulting bias during gradient updates. (arXiv)

Variants and implementation notes

- * Proportional vs rank: proportional uses raw error; rank uses ordering to bound extreme skews. Both trade sample diversity against learning focus. (arXiv)
- * Per-trajectory vs per-move: chess uses full game records. Prioritizing by move-level TD-error concentrates updates on critical positions but may overweight rare tactical outliers. Practical systems (AlphaZero forks, MuZero variants) often keep per-move priorities or per-node visit counts

- * when available. (arXiv)
- * IS correction: necessary as priorities introduce bias. Use an annealed beta that approaches 1 over training. Without IS correction the learned policy/value can diverge from the intended objective. (arXiv)

Empirical notes

- * PER improved sample efficiency in DQN style domains. In game engines, authors of MuZero and follow-ups note PER and related replay heuristics are helpful in practice for faster convergence, though engineering details matter (buffer size, priority exponent, IS schedule). (arXiv)

3. Non-IID sampling strategies beyond PER

Common strategies used or explored in chess/self-play settings:

1. Uniform random from large FIFO buffer. Simple baseline. Reduces short-term correlation but preserves older styles and openings. Used in AlphaZero forks. (arXiv)
2. Reservoir / age-based sampling. Ensures older games remain in buffer at controlled frequency. Helps retain rare endgame motifs. Discussed in reproducibility reports (ELF/AlphaZero reimplementations). (yuandong-tian.com)
3. Reanalyse / offline target refresh. Recompute MCTS targets for stored games with the current network to produce fresher labels. MuZero explicitly introduced reanalyse to squeeze more signal from older trajectories. This reduces distribution shift between stored trajectories and current policy. (arXiv)
4. Curriculum / difficulty scheduling. Sample earlier-phase positions and simpler endgames more or less often depending on network competence. Useful in practice to avoid overfitting to complex tactics early. Curriculum is a general RL practice applied in game agents. (pettingzoo.farama.org)
5. Novelty / coverage sampling. Boost samples from underrepresented openings, rare material imbalances, or high-uncertainty states (e.g., high network entropy or disagreement across ensemble). This targets coverage and reduces blind spots. Recent memory-efficient RL work highlights novelty maps and coverage buffers as anti-forgetting tools. (arXiv)

Tradeoffs

- * Focusing on high-error samples accelerates correction of big mistakes but risks overfitting to noise or outliers.
- * Overemphasizing novelty increases coverage but can dilute gradients from important tactical positions.
- * Reanalyse reduces stale-label problems but costs compute.

4. Sample weighting functions

Three families matter for chess evaluation networks.

1. Importance-sampling (IS) weights. Used to unbiasedly correct prioritized sampling. Weight for sample i typically proportional to $(1 / (N * P(i)))^\beta$ where $P(i)$ is sampling probability and β anneals to 1. Necessary when you want the empirical gradient to approximate the full-buffer objective. PER formalized this. (arXiv)
2. Loss-based / hardness weighting. Weight proportional to $f(\text{Loss}_i)$ where f can be linear, power, or clipped. This emphasizes hard examples. Recent analyses show naive hardness weighting can change optimization dynamics and sometimes harm generalization; properly scheduling and clipping weights is important. Recent 2024-2025 work revisits importance weighting effects in deep nets and shows interactions with optimization and dataset shift are nontrivial. (AIMS)
3. Target-quality / MCTS-derived weights. For self-play with MCTS, use MCTS visit counts or policy entropy as weights. Examples: weight positions proportional to max visit count to emphasize positions where search strongly preferred a move. This ties sampling directly to the search signal and is used in many practical implementations. (arXiv)

Practical recipes for weighting

- * Use IS correction when priorities are steep. Anneal β from small to 1.
- * Clip max weight to avoid extreme gradient variance.
- * Combine signals: e.g., $\text{priority} = \alpha * \text{TD-error_norm} + (1-\alpha) * \text{normalized_visit_count}$. Tune α .
- * When reanalysing, treat refreshed labels as higher-quality and raise their weights modestly.

5. Empirical lessons from chess / Go reimplementations

- * Reproducibility papers and open reimplementations (ELF OpenGo, Leela Chess/Go communities) report that buffer mechanics and sampling hyperparameters materially change learning curves. Uniform sampling with very large buffers can hurt adaptation speed. Smaller buffers with either PER or reanalyse often achieve faster wall-clock progress but may reduce long-term diversity if not managed. (Proceedings of Machine Learning Research)
- * Engineering matters: priority data structures, efficient updates, and parallelism between self-play workers and learners determine whether PER's theoretical gains show up in practice. Many practitioners prefer simpler heuristics unless operating at scale. (yuandong-tian.com)

6. Open problems and research directions

- * Formalizing sample-weight schedules that optimize both short-term correction and long-term generalization for deterministic perfect-information games.
- * Quantifying when reanalyse is better than prioritizing old high-error samples. Reanalyse trades compute for label quality; tradeoff curves are underexplored. (arXiv)
- * Interaction of importance weighting with modern optimizers and large transformer/conv nets. Recent 2025 analyses show classical importance weighting assumptions can fail in deep nets. More targeted theory and benchmarks are needed. (arXiv)

7. Practical checklist for a chess evaluation pipeline

1. Start with uniform sampling from a large FIFO buffer to ensure coverage. (arXiv)
2. Add per-move priority using normalized TD-error or combined TD + visit count. Use rank-based priority if variance is a concern. (arXiv)
3. Apply IS weights with beta annealed toward 1. Clip weights. (arXiv)
4. Implement reanalyse for a fraction of queued games if compute permits. Mark reanalysed trajectories with higher sample quality and raise weight slightly. (arXiv)
5. Maintain explicit coverage monitoring (openings, endgames). Inject underrepresented buckets to prevent blind spots. (Proceedings of Machine Learning Research)

8. References (select)

- * Schaul, T., Quan, J., Antonoglou, I., & Silver, D. 2015. Prioritized Experience Replay.
- * Silver, D. et al. 2017. Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm (AlphaZero).
- * Schrittwieser, J. et al. 2019. MuZero: Mastering Atari, Go, Chess and Shogi by Planning with a Learned Model.
- * Tian, Y. et al. ELF OpenGo and reproducibility notes.
- * Recent analyses on importance weighting and loss-ordered sampling (2024-2025).

3.2.1 Follow-up: Comparing Naive, Information-Theoretic, and Gradient-Based Weighting

```
\begin{verbatim}
Naive heuristics are cheap and often predictive. Information-theoretic and
gradient methods are more precise but cost much more. Below I compare them
from first principles, give practical formulas, and recommend lightweight
hybrids you can use in an AlphaZero/MuZero-style pipeline.

```

1 - What each family measures (intuitively)

1. Naive difficulty heuristics

- * Aim: capture *obvious* learning value using fast, chess-specific signals.
- * Examples: material imbalance, piece activity, mobility, king safety, hanging pieces, captures/checks/promotions, pawn-structure defects, distance to endgame.
- * Why they help: these signals correlate with positions where policy/value need sharper distinctions or where search produced strong corrections.

2. Information-theoretic methods (ensemble disagreement, entropy)

- * Aim: measure epistemic uncertainty in the network's outputs.
- * Examples: variance across ensemble value heads, policy entropy, KL divergence between networks, disagreement on best move.
- * Why they help: they target positions where the model is unsure and can gain most from new data.

3. Gradient-based weighting

- * Aim: measure the local training signal directly.
- * Examples: per-sample gradient norm, influence function approximations, Fisher information, loss change when upweighting a sample.
- * Why they help: they estimate which examples will move parameters most and thus change predictions.

2 - Concrete scoring functions (cheap formulas you can implement now)

Let s be a raw score (higher = more important). Then convert to sampling probability or weight w .

Naive heuristics (0(1) per position):

- * Material imbalance: $s = |\Delta \text{material}|$ (use standard piece values).
- * Mobility: $s = 1 / (1 + \exp(-(mobility - \mu)/\sigma))$ or simply $s = mobility_norm$.
- * Tactical flag: $s = I(\text{capture or check or promotion}) * depth_factor$.
- * Composite: $s = \alpha * \text{norm}(|\Delta \text{material}|) + \beta * \text{norm}(mobility) + \gamma * I(\text{tactical})$. Normalization: $\text{norm}(x) = (x - \min)/(max - \min)$ computed from a rolling window.

Mapping to weight: $w = \text{clip}(s^p, w_{\min}, w_{\max})$ or softmax over batch: $w_i = \exp(\lambda * s_i) / \sum \exp(\lambda * s_j)$. Use p in $[0.5, 2]$, λ small to avoid extreme skew.

Information-theoretic (cost $\sim k$ forward passes for ensemble of k nets):

- * Ensemble variance (value): $s = \text{Var}_k(v_k(x))$.

- * Policy disagreement: $s = \text{KL}(\pi_{\text{ensemble}} || \pi_{\text{mean}})$ or JS divergence.
- * Entropy: $s = -\sum_a \pi(a) \log \pi(a)$.
Weight mapping same as above. If you use a 3-5 member ensemble run forwards in batch, cost approximately 3-5x inference.

Gradient-based (cost ~1 backward per sample or approximation):

- * Gradient norm: $s = \|\text{grad}_{\theta} L(x)\|$ (per-sample). Exact requires backward pass per sample.
- * Last-layer approx: s approximately $\|\text{grad}_{\theta_{\text{last}}} L(x)\|$ computed cheaply by restricting to last layer.
- * Influence approximation: s approximately $\text{grad } L(x)^T H^{-1} g_{\text{val}}$ where H is approximate Hessian. Use diagonal or Fisher approximation.
Map to weight via clipping and annealing as above.

3 - Computational cost (big-O style) and latency

- * Naive heuristics: $O(1)$ per position. Adds negligible overhead to replay sampling or generation. Scales to millions of positions.
- * Ensemble/entropy: $O(k*F)$ where F is forward cost. For $k=3$ and modern nets F is significant. Can be done offline or sampled fractionally. Memory and IO impact if stored with replay.
- * Per-sample gradient norm: $O(B*F + N*B_{\text{grad}})$ depending on how you compute.
Exact per-sample backward is roughly equal to a full training pass per sample and is usually infeasible. Last-layer gradient or batch-level proxies reduce cost by 1-2 orders. Influence methods need Hessian approximations and are heavy.

Practical numbers (ballpark):

- * Naive: <1 microsecond to 10 microseconds per position (feature arithmetic).
- * Ensemble 3 models: ~3x forward time. If forward is 1ms, ensemble ~3ms.
- * Per-sample full backward: ~1-2x forward per sample but multiplied by batch size; effectively too slow if done for each candidate sample in large buffers.

4 - Correlation with true learning value and failure modes

- * Naive heuristics
 - * Correlation: moderate. Good at catching tactical and material signals.
 - * Failures: misses subtle strategic confusion. Can overweight easy but obvious positions. Sensitive to hand-tuned thresholds. Vulnerable to adversarial or distribution shift where heuristics misalign with network errors.
- * Ensemble/entropy

- * Correlation: higher with epistemic uncertainty. Targets positions where model lacks knowledge.
- * Failures: expensive. Can be fooled by ensembles that share biases. Requires diverse ensemble or dropout schedules. High entropy can occur in trivially drawn positions with many equivalent moves.

* Gradient-based

- * Correlation: highest in principle. Directly measures how much a parameter update would change loss.
- * Failures: noisy, can be dominated by outliers and label noise. Computationally heavy. Influence approximations break when curvature assumptions fail.

5 - Practical tradeoffs and engineering patterns

1. Use cheap heuristics as primary filter. They are cheap and reduce search space for expensive methods.

Example: mark positions with tactical flag or large material change and compute ensemble/gradient scores only for those.

2. Batch expensive computations and cache results.

Compute ensemble disagreement once per trajectory generation or during reanalyse. Store a small metadata vector with the replay entry.

3. Last-layer gradient as a reasonable compromise.

Cost: one backward restricted to last layer. Correlates with full gradient but far cheaper.

4. Use rank-based priority, not raw weights.

Rank reduces sensitivity to scale and outliers. Implementation: maintain priority tree keyed by rank or use reservoir buckets.

5. Cap and anneal.

Clip weights to avoid gradient explosion. Anneal emphasis from high to low to encourage stability.

6. Combine signals via convex mix.

```
priority_score = alpha*norm(heuristic) + beta*norm(ensemble_var) +
gamma*norm(grad_proxy). Tune alpha, beta, gamma. Default alpha=0.6,
```

beta=0.3, gamma=0.1 is a practical starting point.

7. Importance-sampling correction.

If sampling non-uniformly, apply IS weights during gradient update: multiply per-sample gradient by $(1 / P(i))^\beta$ where beta approaches 1 over training.

6 - Recommendations for an intern experiment plan

1. Baseline: uniform sampling from large FIFO buffer. Measure learning curve.
2. Add naive heuristic prioritization (composite of tactical flag + material imbalance). Use power weighting p=1. Clip weights to [0.5, 2.0]. Measure wall-clock and sample efficiency.
3. Add a 3-model ensemble disagreement computed only for samples with top-10% heuristic score. Use ensemble variance to re-rank that slice. Store ensemble score in buffer metadata. Compare.
4. Replace ensemble with last-layer gradient norm on that same slice. Compare cost and performance.
5. Ablate annealing and clipping. Track instability, overfitting, and long-term generalization (cross-validated on held-out openings and endgames).

Metrics: Elo vs holdout positions, calibration of value head, policy top-1 accuracy, and catastrophic forgetting measured by coverage loss over openings.

7 - Short prescriptive summary

- * If compute is constrained: use naive heuristics + rank-based priority + IS correction. Cheap and robust.
- * If moderate compute available: add last-layer gradient norm or small ensemble on a filtered subset. Best tradeoff between signal and cost.
- * If compute abundant and you want maximal sample efficiency: full ensemble + periodic per-sample gradient/influence estimates with careful clipping and caching. Expect diminishing returns and engineering complexity.

If you want, I will:

- * produce code snippets for computing these scores (feature extraction, rank priority tree, last-layer gradient hook) suitable for PyTorch.

3.3 Observations on Quality and Accuracy